

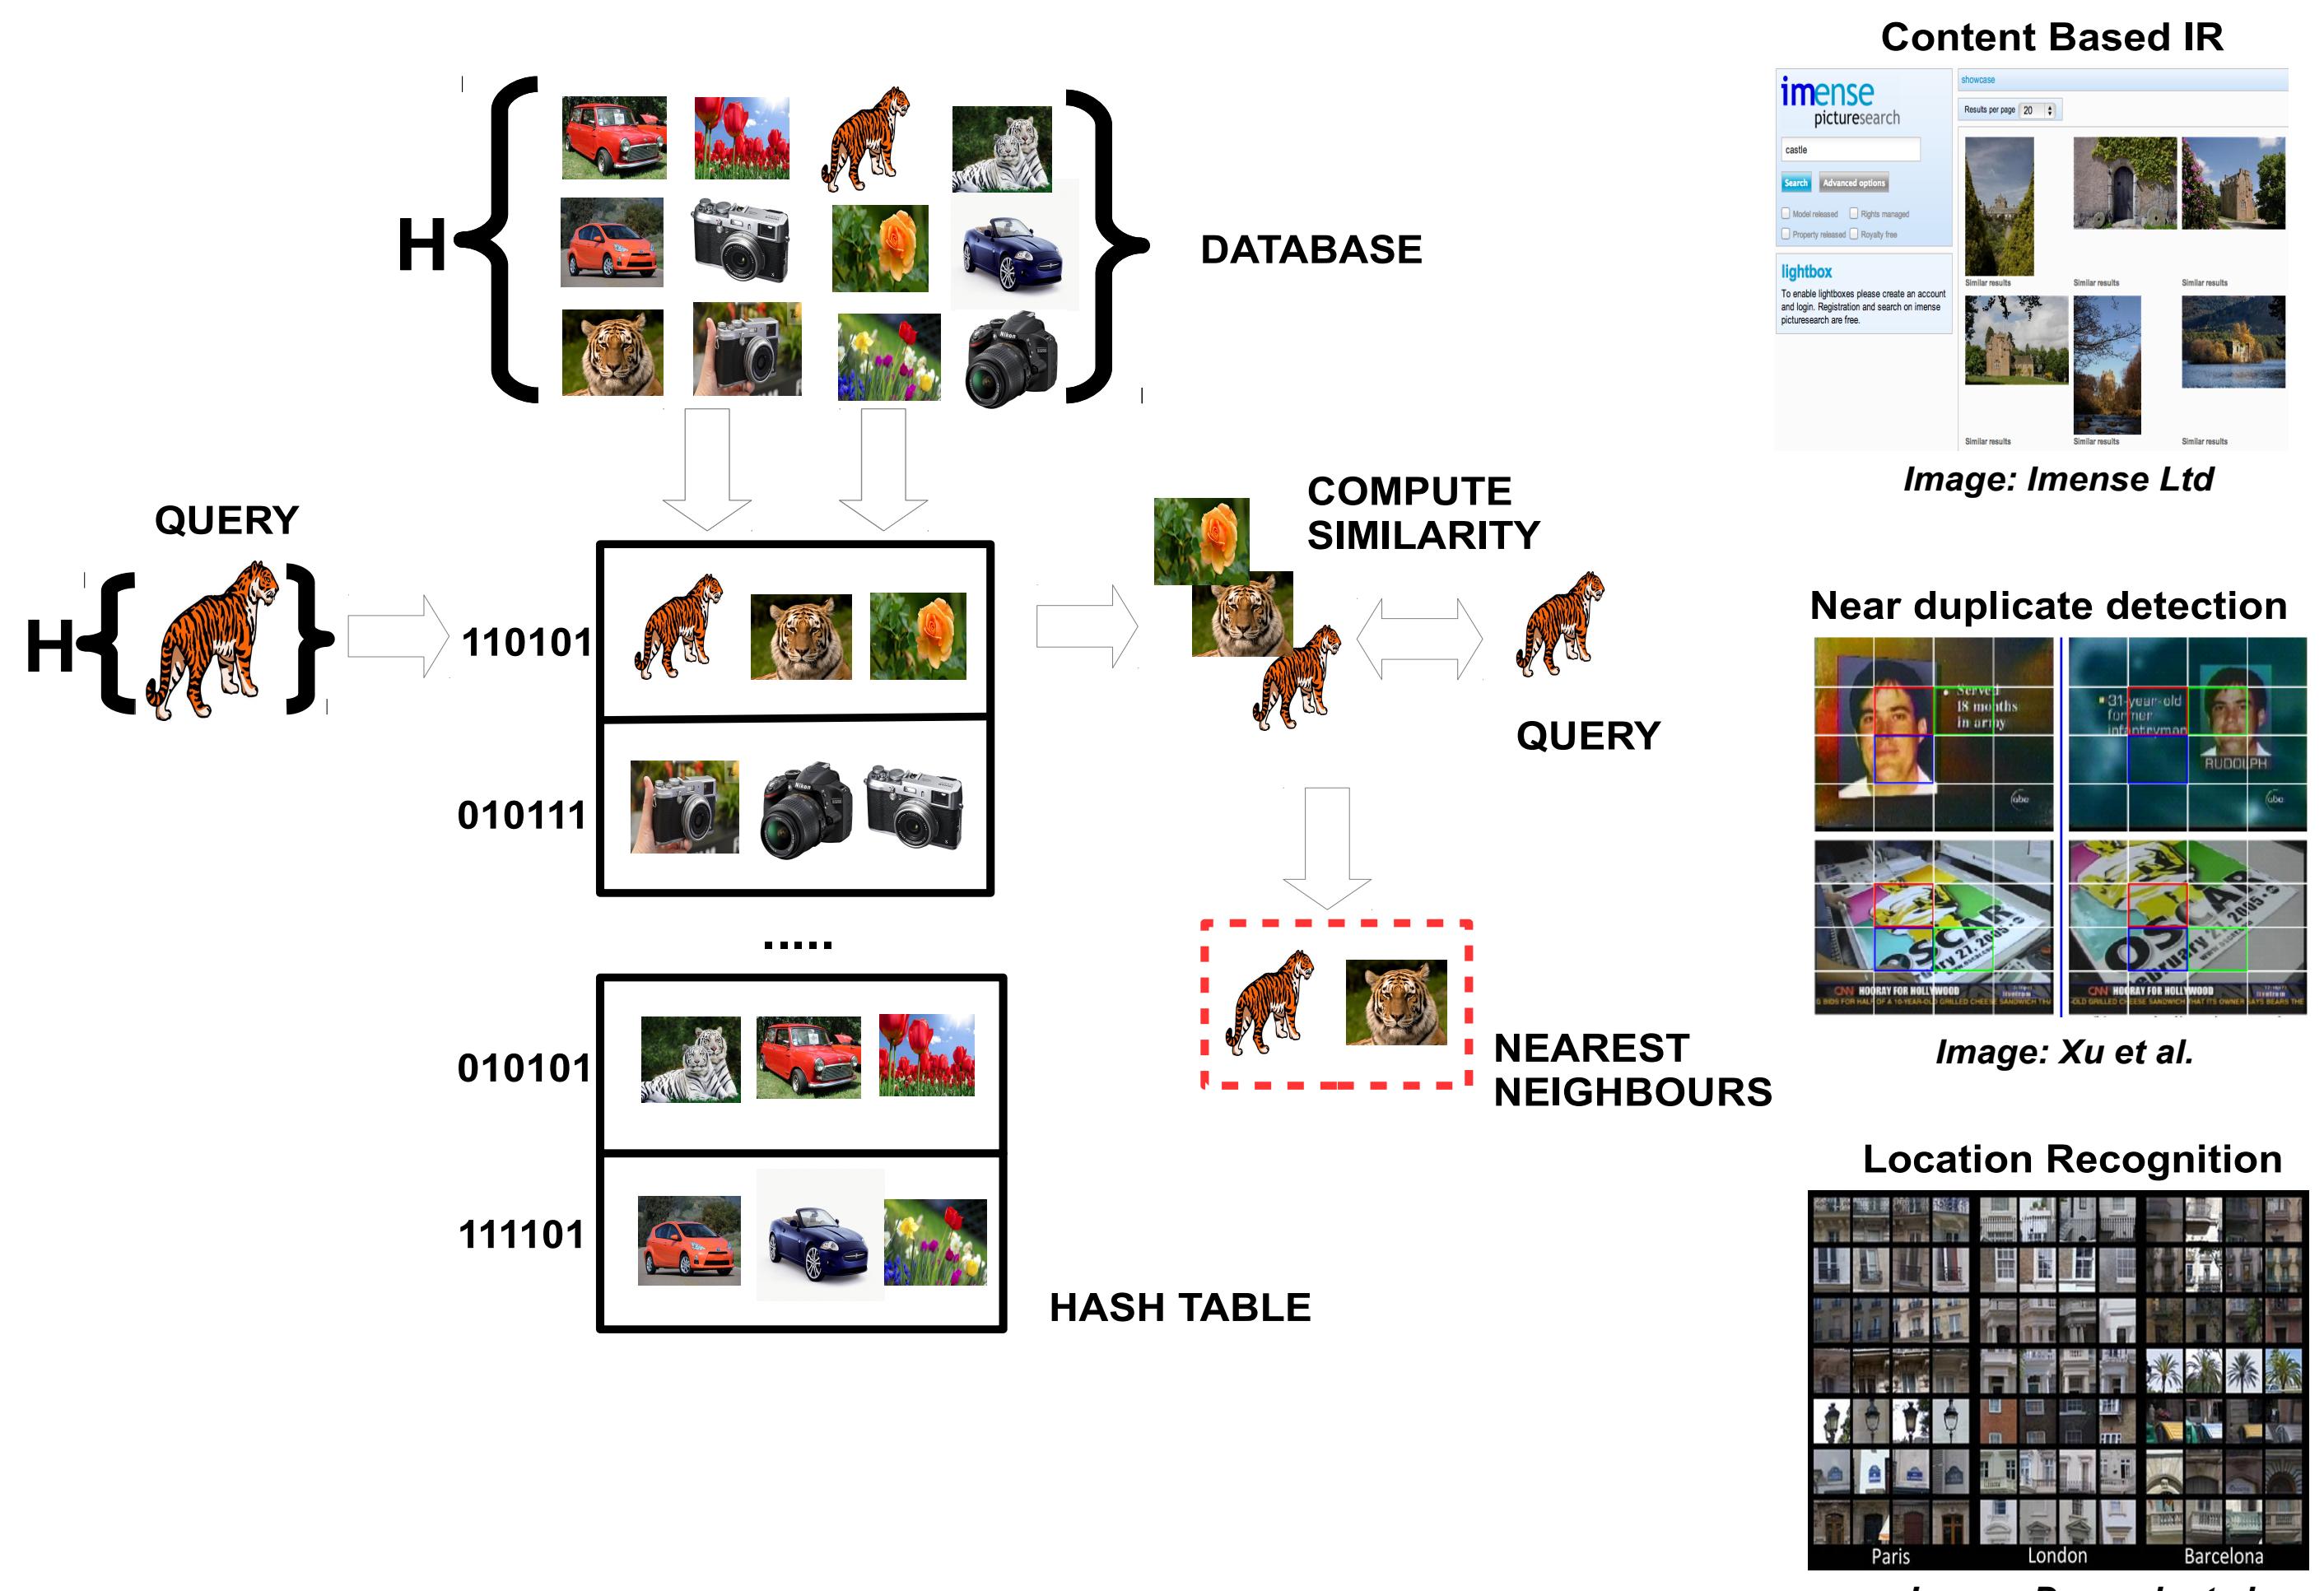


RESEARCH QUESTION

- Locality sensitive hashing (LSH) [1] fractures the input feature space space with randomly placed hyperplanes.
- Can we do better by using supervision to adjust the hyperplanes?

INTRODUCTION

- Problem:** Constant time nearest-neighbour search in large datasets.
- Hashing-based approximate nearest neighbour (NN) search:**
 - Index image/documents into the buckets of hashtable(s)
 - Encourage collisions between similar images/documents



Advantages:

- $O(1)$ lookup per query rather than $O(N)$ (brute-force)
- Memory/storage saving due to compact binary codes

GRAPH REGULARISED HASHING (GRH)

- We propose a two step *iterative* hashing model, Graph Regularised Hashing (GRH) [5]. GRH uses supervision in the form of an adjacency matrix that specifies whether or not data-points are related.

- Step A: Graph Regularisation:** the K-bit hashcode of a data-point is set to the average of the data-points of its nearest neighbours as specified by the adjacency graph:

$$\mathbf{L}_m \leftarrow \text{sgn}(\alpha \mathbf{SD}^{-1} \mathbf{L}_{m-1} + (1-\alpha) \mathbf{L}_0)$$

- * \mathbf{S} : Affinity (adjacency) matrix
- * \mathbf{D} : Diagonal degree matrix
- * \mathbf{L} : Binary bits at iteration m
- * $\alpha \in \{0, 1\}$: Linear interpolation parameter

- Step A is a simple sparse-sparse matrix multiplication, and can be implemented very efficiently. Any existing hash function e.g. LSH [1] can be used to initialise the bits in \mathbf{L}_0

- **Step B: Data-Space Partitioning:** the hashcodes produced in Step A are used as the *labels* to learn K binary classifiers. This is the out-of-sample extension step, allowing the encoding of data-points not seen before:

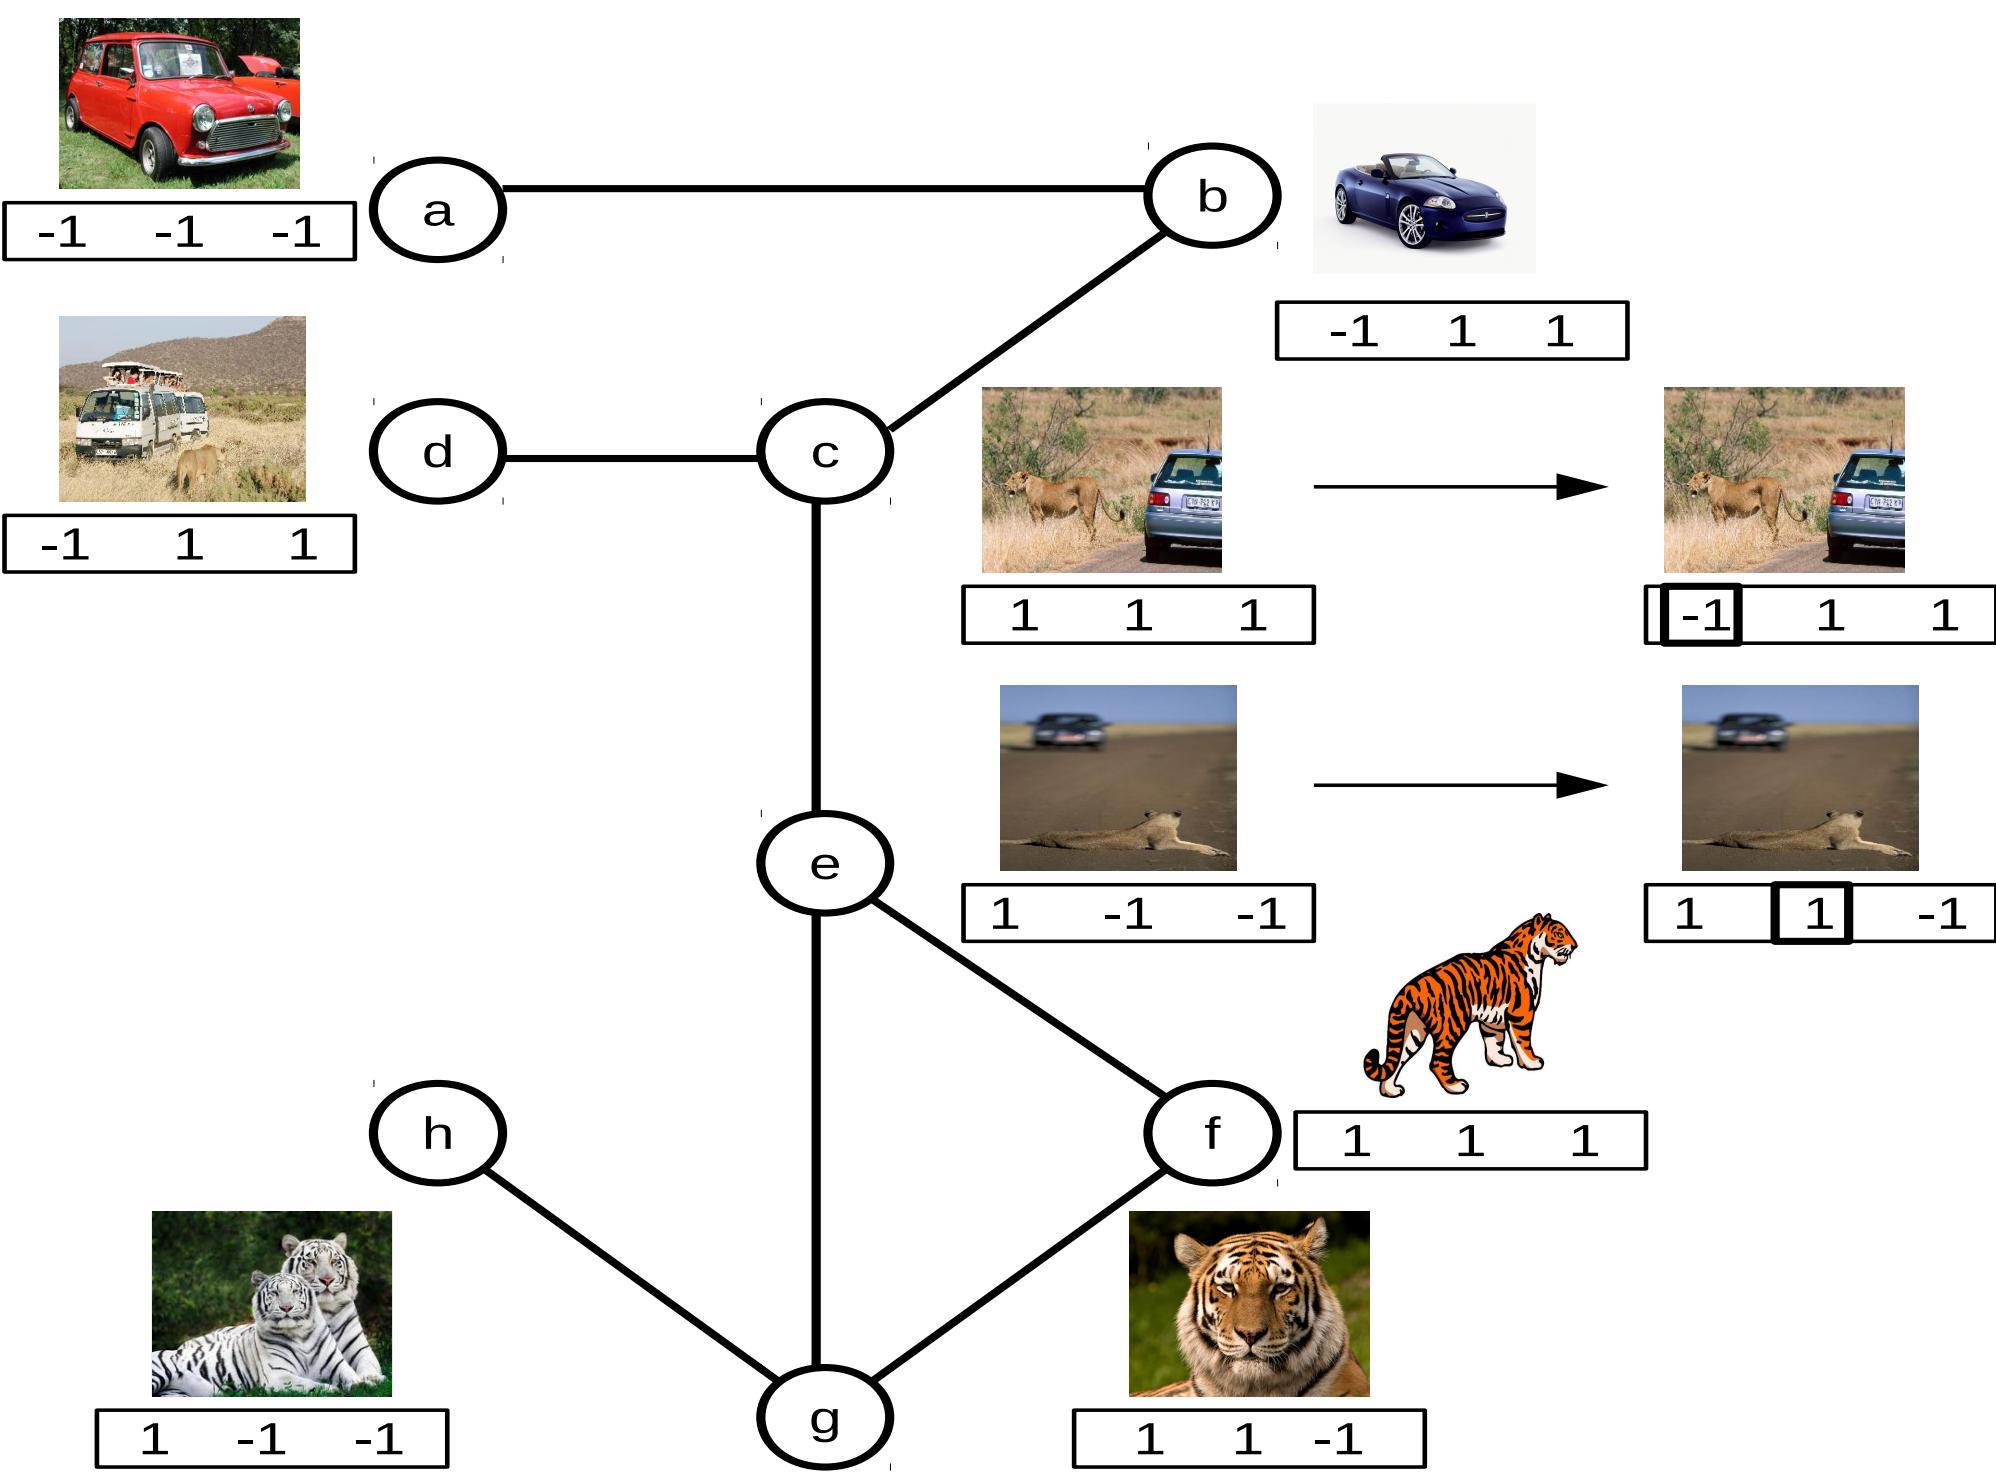
$$\begin{aligned} \text{for } k = 1 \dots K : \quad & \min \| \mathbf{w}_k \|^2 + C \sum_{i=1}^{N_{trd}} \xi_{ik} \\ \text{s.t.} \quad & L_{ik} (\mathbf{w}_k^\top \mathbf{x}_i + t_k) \geq 1 - \xi_{ik} \quad \text{for } i = 1 \dots N_{trd} \end{aligned}$$

- * \mathbf{w}_k : Hyperplane k t_k : bias k
- * \mathbf{x}_i : data-point i L_{ik} : bit k of data-point i
- * ξ_{ik} : slack variable K : # bits N_{trd} : # data-points

- Steps A-B are repeated for a set number of iterations (M) (e.g. < 10). The learnt hyperplanes \mathbf{w}_k can then be used to encode unseen data-points (via a simple dot-product).

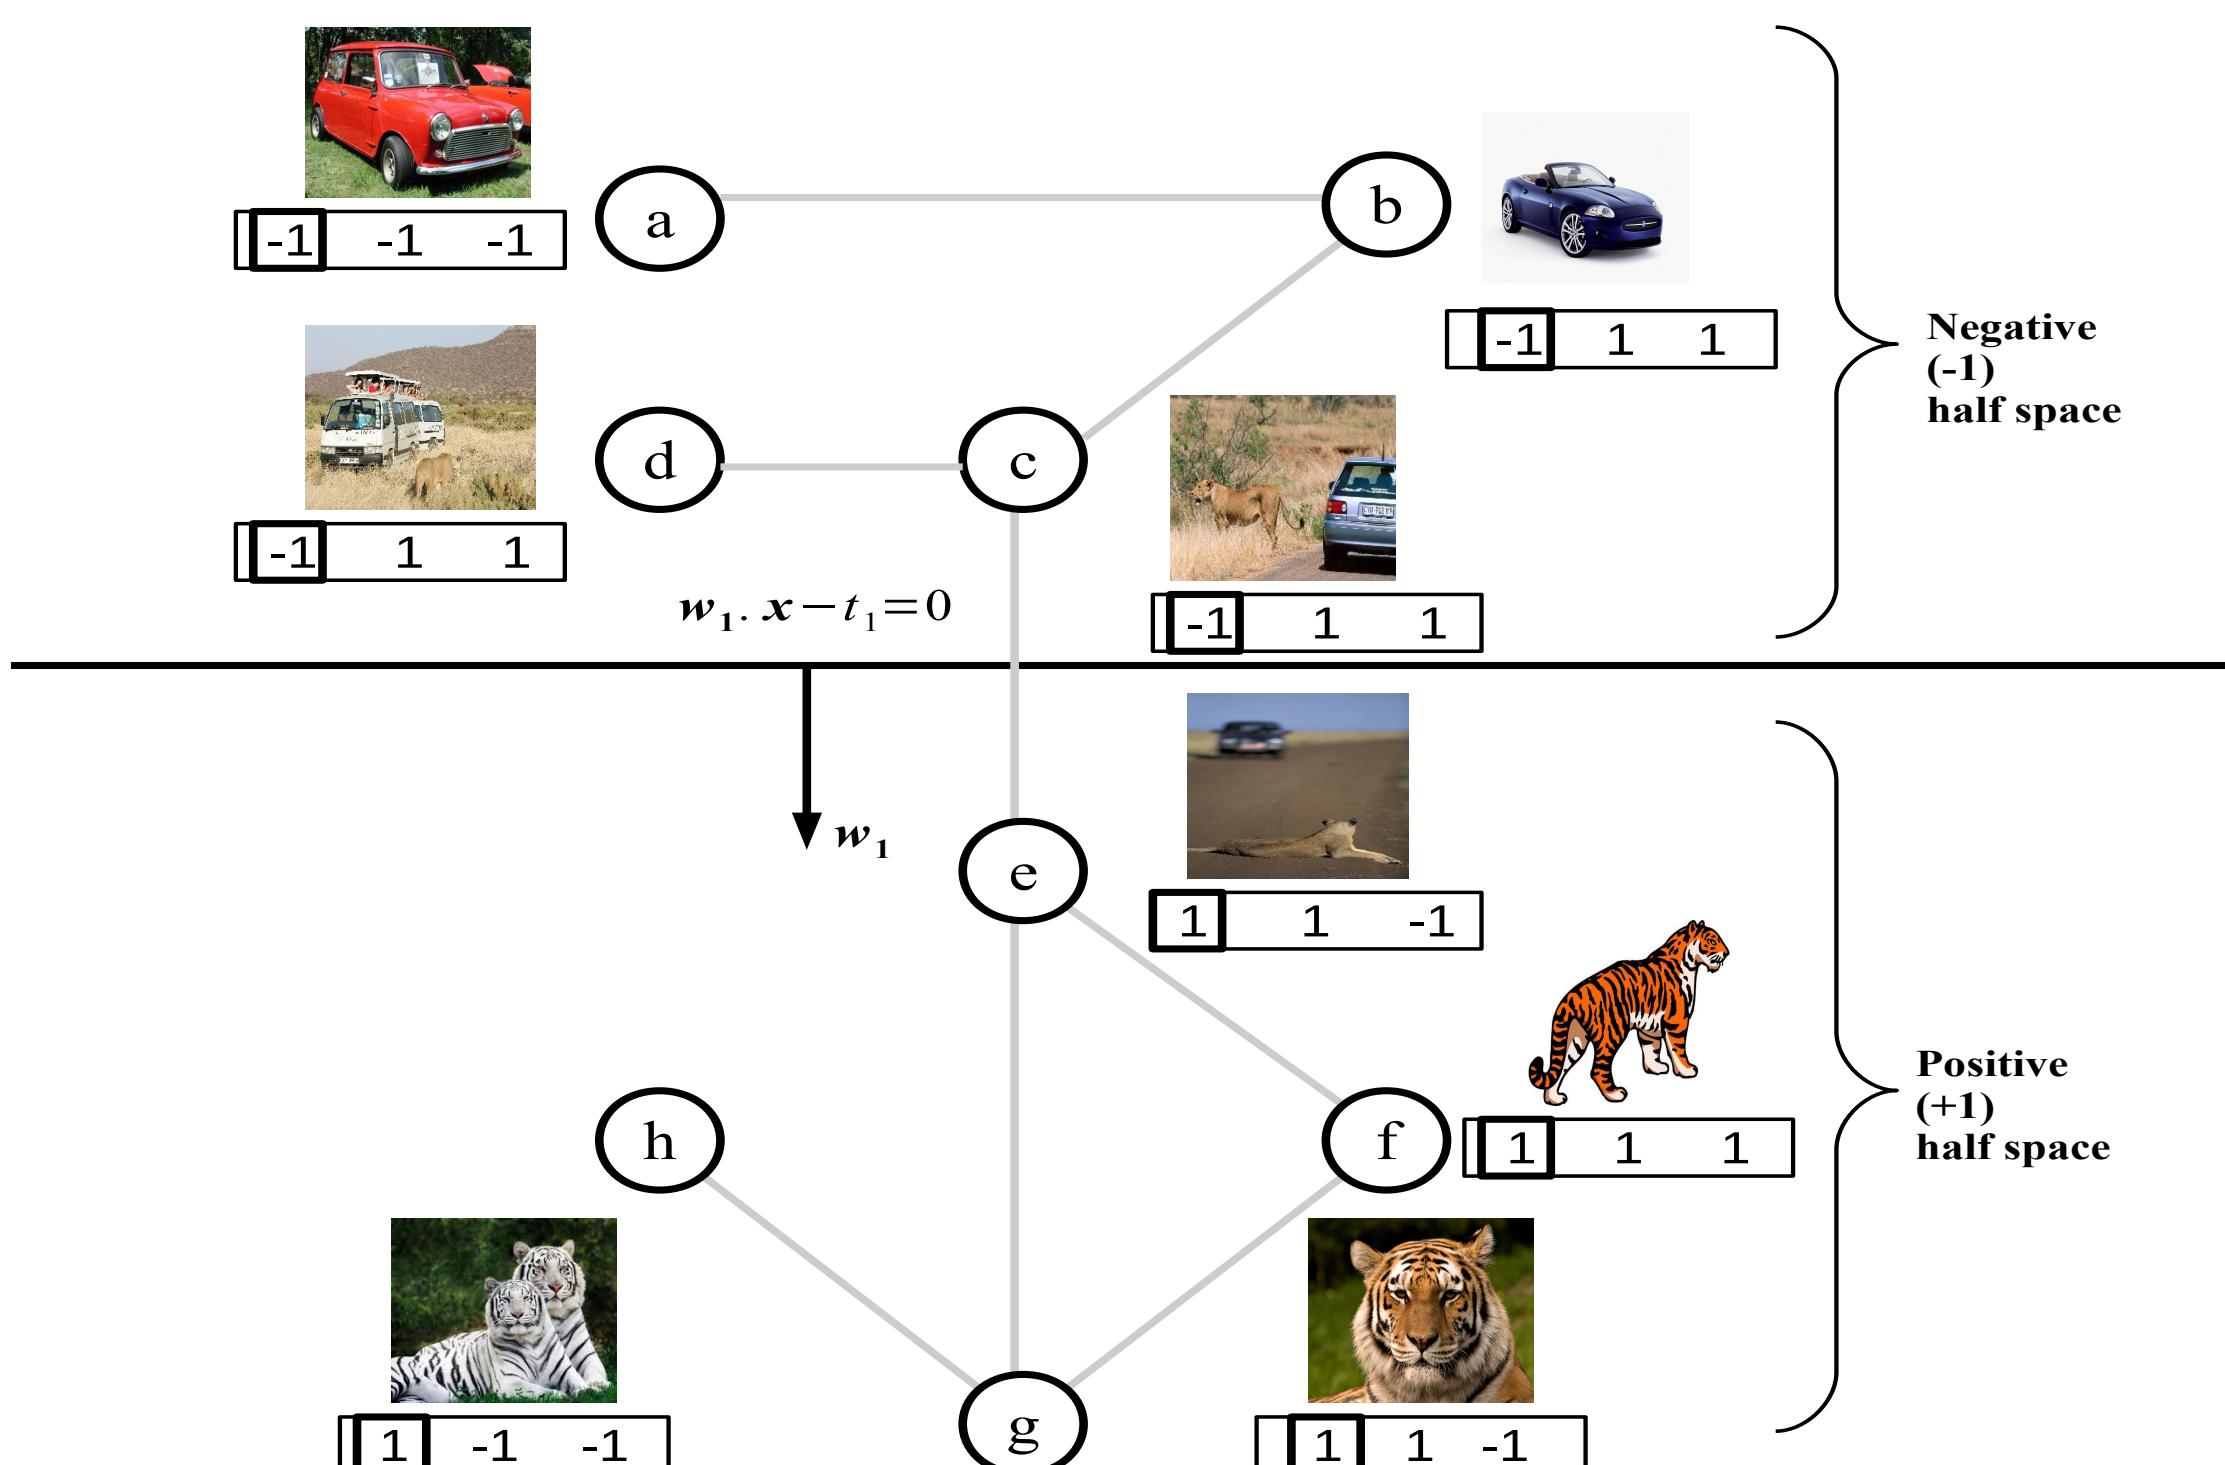
STEP A: GRAPH REGULARISATION

- Toy example: nodes are images with 3-bit LSH encoding. Arcs indicate nearest neighbour relationships. We show two images (c,e) having their hashcodes updated in Step A:



STEP B: DATA-SPACE PARTITIONING

- Here we show a hyperplane being learnt using the first bit as (highlighted with bold box) as label. One hyperplane is learnt per bit.



QUANTITATIVE RESULTS (CIFAR-10) (MORE DATASETS IN PAPER)

- Mean average precision (mAP) image retrieval results using GIST features on CIFAR-10 ($\blacktriangle\blacktriangle$: is significant at $p < 0.01$):

	CIFAR-10			
	16 bits	32 bits	48 bits	64 bits
ITQ+CCA [2]	0.2015	0.2130	0.2208	0.2237
STH [3]	0.2352	0.2072	0.2118	0.2000
KSH [4]	0.2496	0.2785	0.2849	0.2905
GRH [5]	0.2991$\blacktriangle\blacktriangle$	0.3122$\blacktriangle\blacktriangle$	0.3252$\blacktriangle\blacktriangle$	0.3350$\blacktriangle\blacktriangle$

- Timings (seconds) averaged over 10 runs. GRH is 1) faster to train and 2) is faster to encode unseen data-points:

	Training	Testing	Total
GRH [5]	8.01	0.03	8.04
KSH [4]	74.02	0.10	74.12
BRE [6]	227.84	0.37	228.21

SUMMARY OF KEY FINDINGS

- First both accurate and scalable supervised hashing model
- Future work will extend GRH to streaming data sources
- Code online: <https://github.com/sjmoran/grh>

References:

- P. Indyk, R. Motwani: Approximate nearest neighbors: Towards removing the curse of dimensionality. In: STOC (1998).
- Y. Gong, S. Lazebnik: Iterative Quantisation. In: CVPR (2011). , [3] D. Zhang et al. Self-Taught Hashing. In: SIGIR (2010)., [4] W. Liu et al. Supervised Hashing with Kernels. In: CVPR (2012)., [5] S. Moran, V. Lavrenko. Graph Regularised Hashing. In: ECIR (2015).,
- B. Kulis, T. Darrell et al. Binary Reconstructive Embedding. In: NIPS (2009).